Parametric Models Fitting with Automatic Differentiation

Gabriel Peyré

www.numerical-tours.com

ENS
ÉCOLE NORMALE SUPÉRIEURE

Mathematical Coffees

Huawei-FSMP joint seminars

https://mathematical-coffees.github.io

HUAWEI

FSMP Fondation Sciences Mathématiques de Paris

Organized by: Mérouane Debbah & Gabriel Peyré

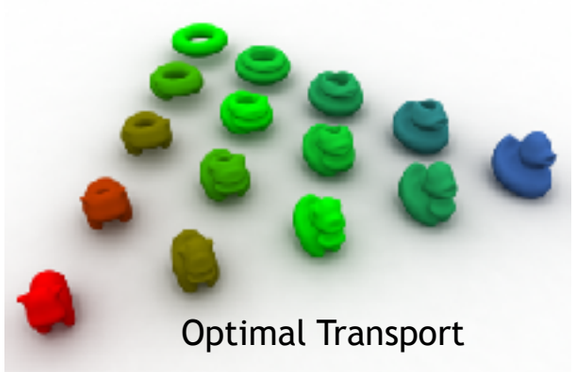

Optimal Transport

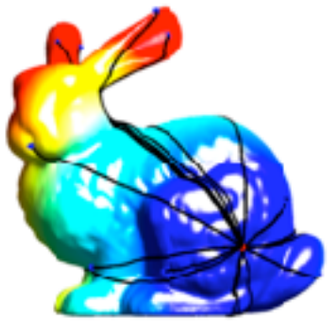

Geodesics

Meshes

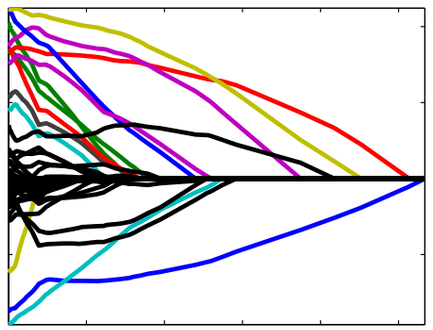

Optimization

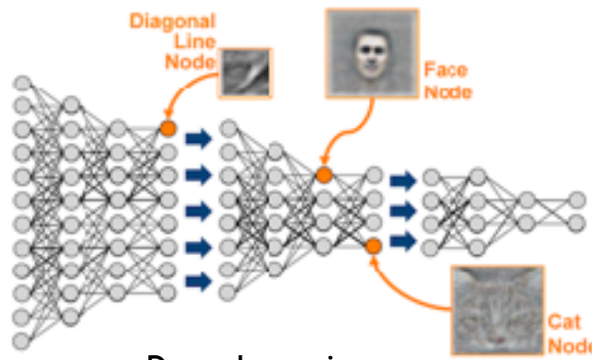

Deep Learning

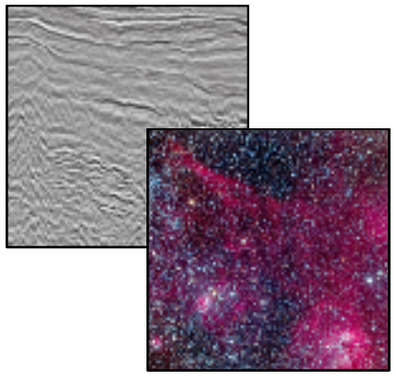

Sparsity

Neuro-imaging

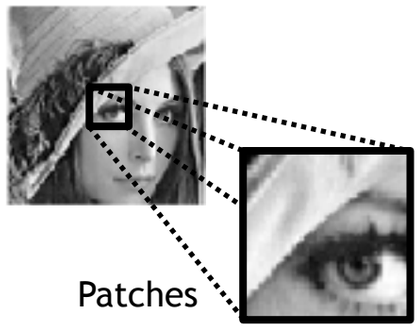

Patches

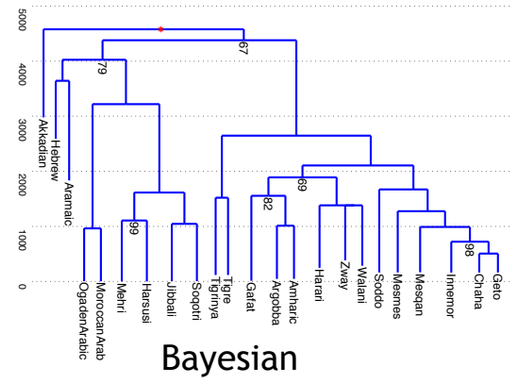

Bayesian

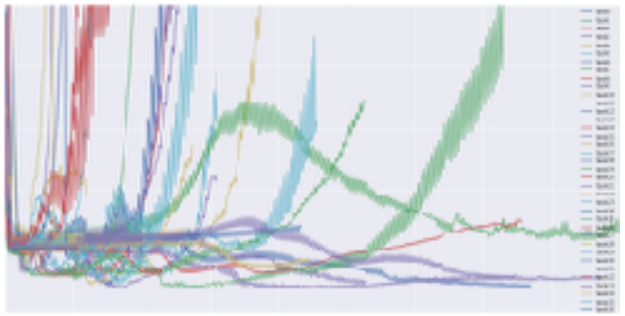

Parallel/Stochastic

Alexandre Allauzen, Paris-Sud.
Pierre Alliez, INRIA.
Guillaume Charpiat, INRIA.
Emilie Chouzenoux, Paris-Est.

Nicolas Courty, IRISA.
Laurent Cohen, CNRS Dauphine.
Marco Cuturi, ENSAE.
Julie Delon, Paris 5.

Fabian Pedregosa, INRIA.
Julien Tierny, CNRS and P6.
Robin Ryder, Paris-Dauphine.
Gael Varoquaux, INRIA.

Jalal Fadili, ENSICaen.
Alexandre Gramfort, INRIA.
Matthieu Kowalski, Supelec.
Jean-Marie Mirebeau, CNRS,P-Sud.

Parametric Models

(Noisy) observations (x_i, y_j), try to infer $y = f(x)$.

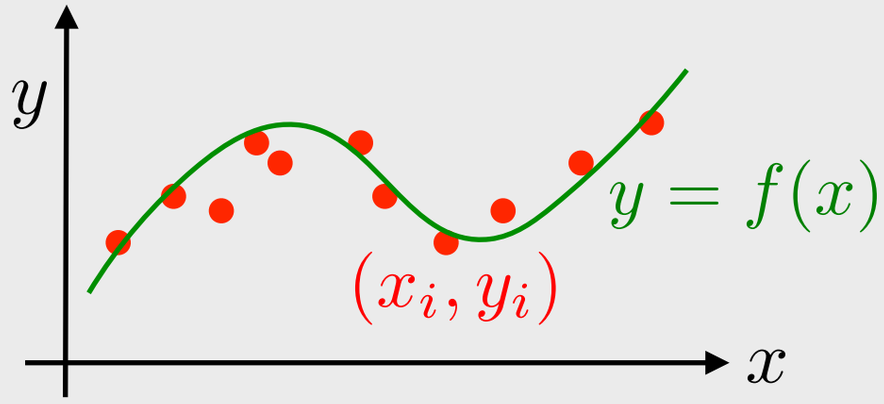

Regression $(x, y) \in \mathbb{R}^n \times \mathbb{R}^p$

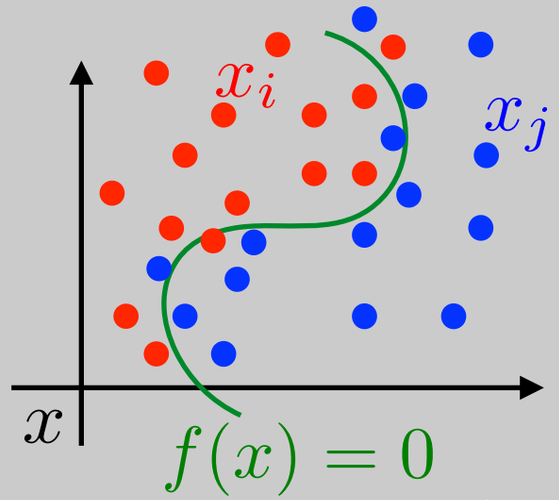

$f_i = -1$

$f_j = 1$

Classification $(x, y) \in \mathbb{R}^n \times \{-1, 1\}$

Parametric Models

(Noisy) observations (x_i, y_j), try to infer $y = f(x)$.

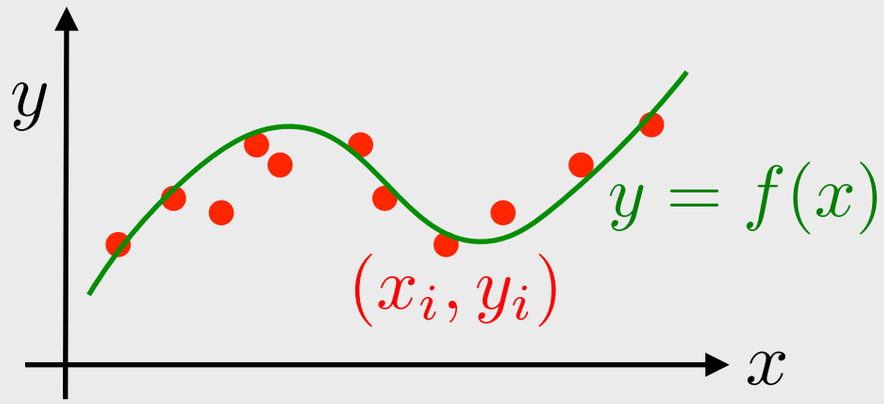

Regression $(x, y) \in \mathbb{R}^n \times \mathbb{R}^p$

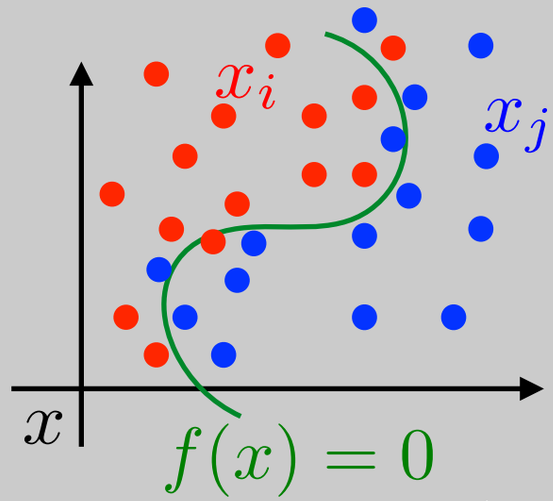

$f_i = -1$

$f_j = 1$

Classification $(x, y) \in \mathbb{R}^n \times \{-1, 1\}$

Parametric model: $y = f(x, \theta)$, find θ.

Linear model: $f(x, \theta) = \langle x, \theta \rangle$.

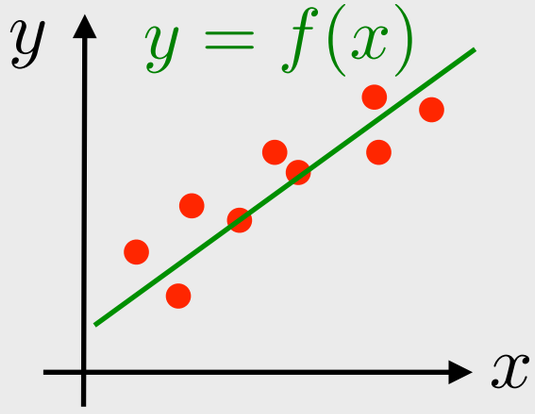

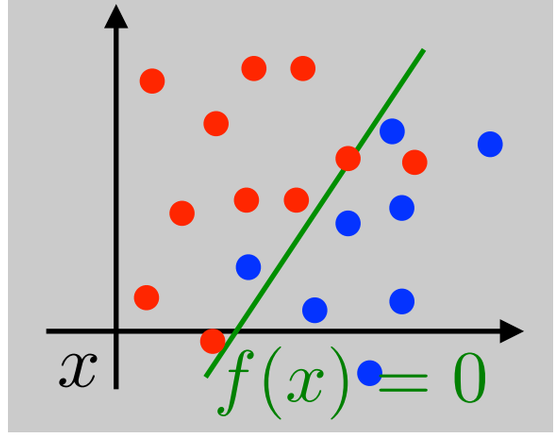

Parametric Models

(Noisy) observations (x_i, y_j), try to infer $y = f(x)$.

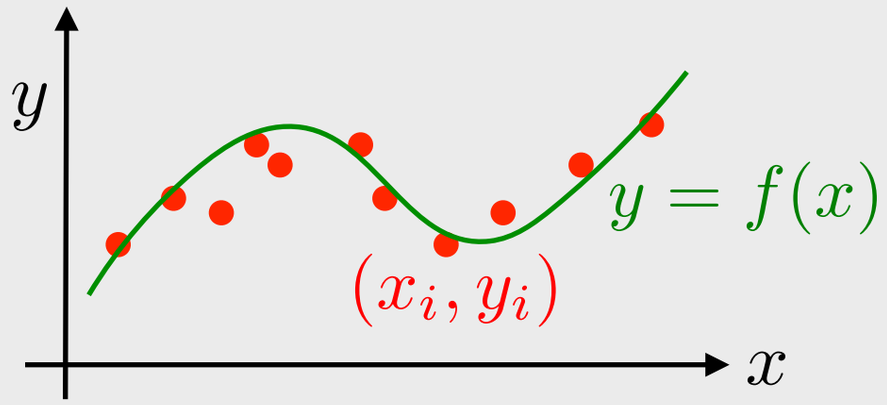

Regression $(x, y) \in \mathbb{R}^n \times \mathbb{R}^p$

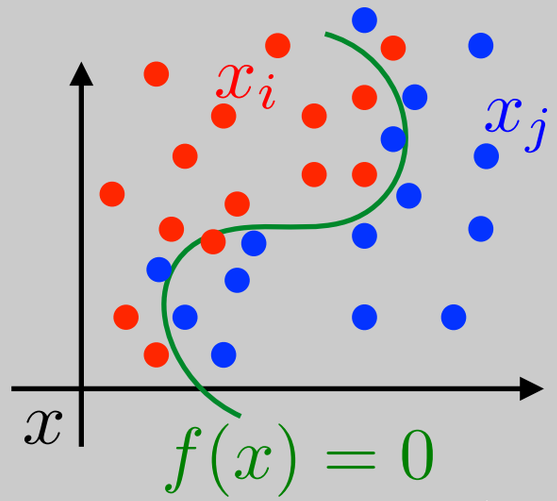

$f_i = -1$

$f_j = 1$

Classification $(x, y) \in \mathbb{R}^n \times \{-1, 1\}$

Parametric model: $y = f(x, \theta)$, find θ.

Linear model: $f(x, \theta) = \langle x, \theta \rangle$.

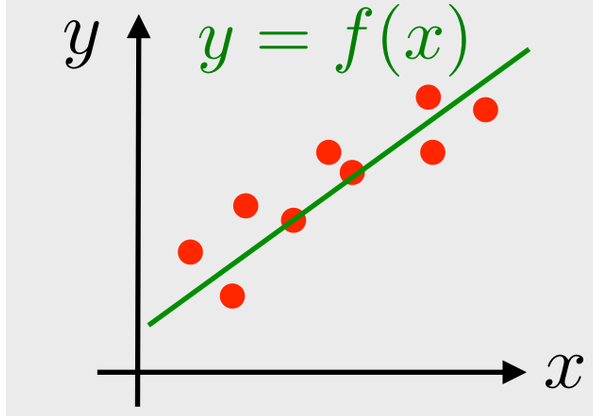

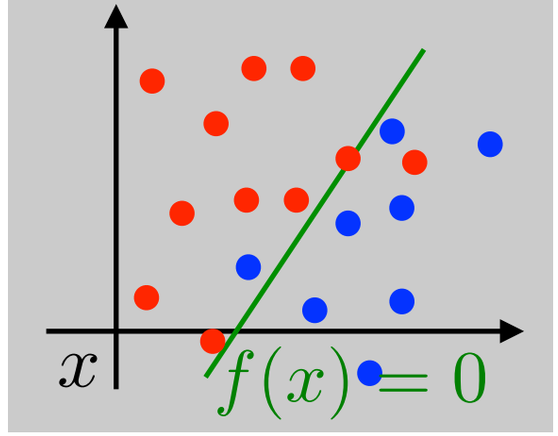

Deep network:

$$f(x, \theta) = \theta_K(\ldots \rho(\theta_2(\rho(\theta_1(x)\ldots)$$

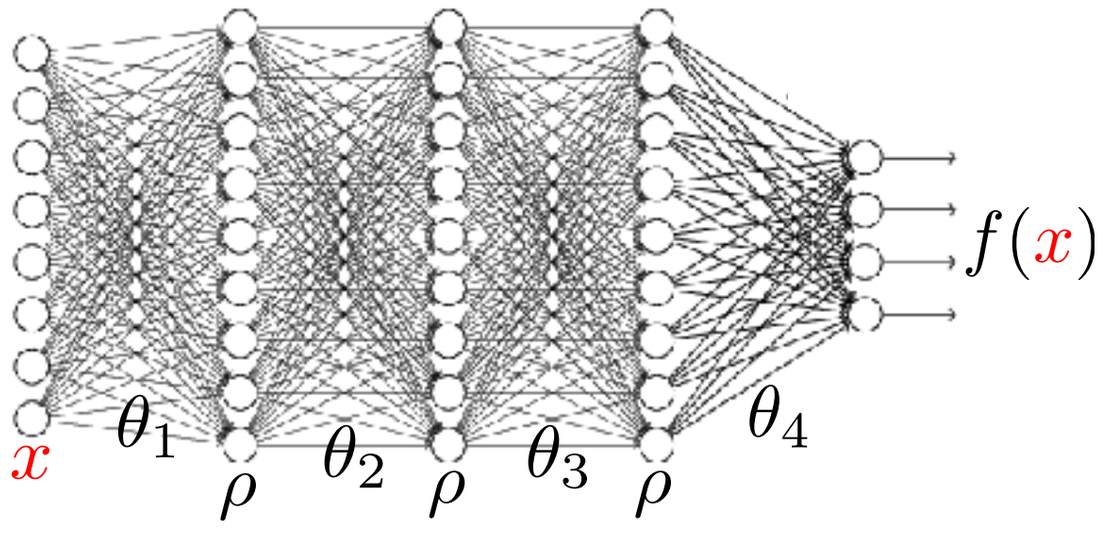

Empirical Loss Minimization

Regression: $(y, y') \in \mathbb{R}^d \times \mathbb{R}^d, \qquad L(y, y') = \|y - y'\|^2$

Classification: $(y, y') \in \mathbb{R}^d \times \{-1, 1\}, \qquad L(y, y') = \log(\exp(-y'y) + 1)$

Loss minimization:

$$\min_\theta \sum_i L(f(x_i, \theta), y_i)$$

$$\min_\theta \mathbb{E}_{(X,Y)}(L(f(X, \theta), Y))$$

Empirical Loss Minimization

Regression: $(y, y') \in \mathbb{R}^d \times \mathbb{R}^d, \qquad L(y, y') = \|y - y'\|^2$

Classification: $(y, y') \in \mathbb{R}^d \times \{-1, 1\}, \qquad L(y, y') = \log(\exp(-y'y) + 1)$

Loss minimization:

$$\min_\theta \sum_i L(f(x_i, \theta), y_i) \qquad \min_\theta \mathbb{E}_{(X,Y)}(L(f(X, \theta), Y))$$

Stochastic gradient descent:

– Sample: $(x, y) \in \{(x_i, y_i)\}_i \qquad (x, y) \sim (X, Y)$

– Update: $\theta^{(\ell+1)} \stackrel{\text{def.}}{=} \theta^{(\ell)} - \tau_\ell \nabla_\theta \ell_{x,y}(\theta)$

where $\ell_{x,y}(\theta) \stackrel{\text{def.}}{=} L(f(x, \theta), y)$

Automatic Differentiation

How to compute $\nabla \ell_{x,y}(\theta)$? $\quad\quad$ $\ell_{x,y}(\theta) \stackrel{\text{def.}}{=} L(f(x,\theta),y)$

Chain rule: $\nabla \ell_{x,y}(\theta) = [\partial f(x,\theta)]^\top (\nabla L(f(x,\theta),y))$

Linear $f(x,\theta) = \theta \times x$: $\partial f(x,\theta) = \theta$.

Non-linear $f(x,\theta)$: painful ... but $\ell_{x,y}$ it is just a computer program.

Automatic Differentiation

How to compute $\nabla \ell_{x,y}(\theta)$? $\qquad \ell_{x,y}(\theta) \stackrel{\text{def.}}{=} L(f(x,\theta),y)$

Chain rule: $\nabla \ell_{x,y}(\theta) = [\partial f(x,\theta)]^\top (\nabla L(f(x,\theta),y))$

Linear $f(x,\theta) = \theta \times x$: $\partial f(x,\theta) = \theta$.

Non-linear $f(x,\theta)$: painful ... but $\ell_{x,y}$ it is just a computer program.

Computer program $\Leftrightarrow$ directed acyclic graph $\Leftrightarrow$ linear ordering of nodes $(\theta_r)_r$

forward

```
function ℓ(θ_1, ..., θ_M)
    for r = M+1, ..., R
        θ_r = g_r(θ_Parents(r))
    return θ_R
```

computing ℓ

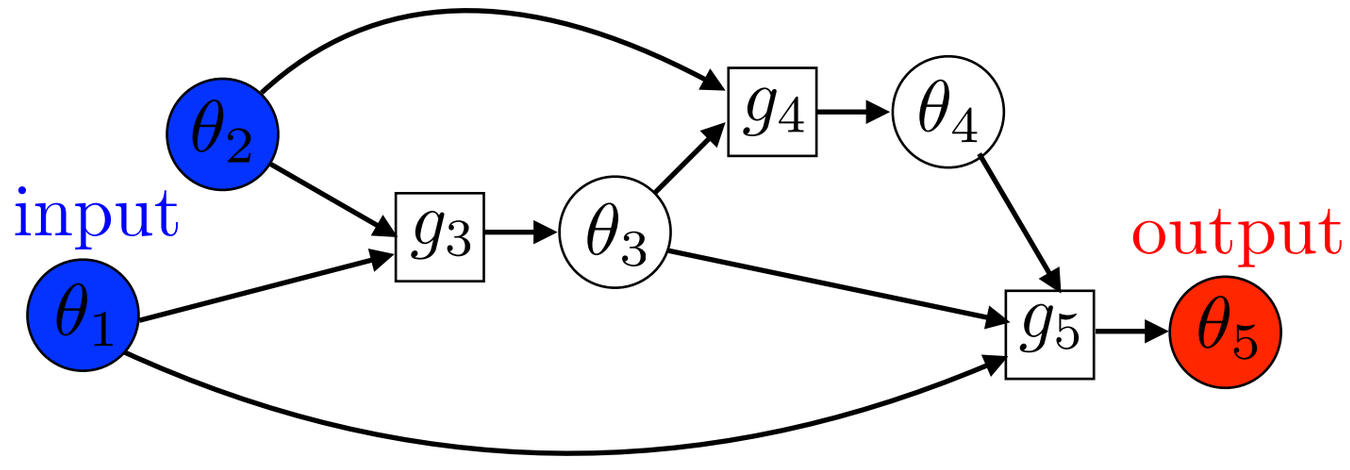

Automatic Differentiation

How to compute $\nabla \ell_{x,y}(\theta)$? $\qquad \ell_{x,y}(\theta) \stackrel{\text{def.}}{=} L(f(x,\theta),y)$

Chain rule: $\nabla \ell_{x,y}(\theta) = [\partial f(x,\theta)]^\top (\nabla L(f(x,\theta),y))$

Linear $f(x,\theta) = \theta \times x$: $\partial f(x,\theta) = \theta$.

Non-linear $f(x,\theta)$: painful ... but $\ell_{x,y}$ it is just a computer program.

Computer program $\Leftrightarrow$ directed acyclic graph $\Leftrightarrow$ linear ordering of nodes $(\theta_r)_r$

forward

```
function ℓ(θ_1, ..., θ_M)
    for r = M+1, ..., R
        θ_r = g_r(θ_Parents(r))
    return θ_R
```

backward

```
function ∇ℓ(θ_1, ..., θ_M)
    ∇_R ℓ = 1
    for r = R−1, ..., 1
        ∇_r ℓ = Σ_{s∈Child(r)} ∂_r g_s(θ) ∇_s ℓ
    return (∇_1 ℓ, ..., ∇_M ℓ)
```

computing ℓ

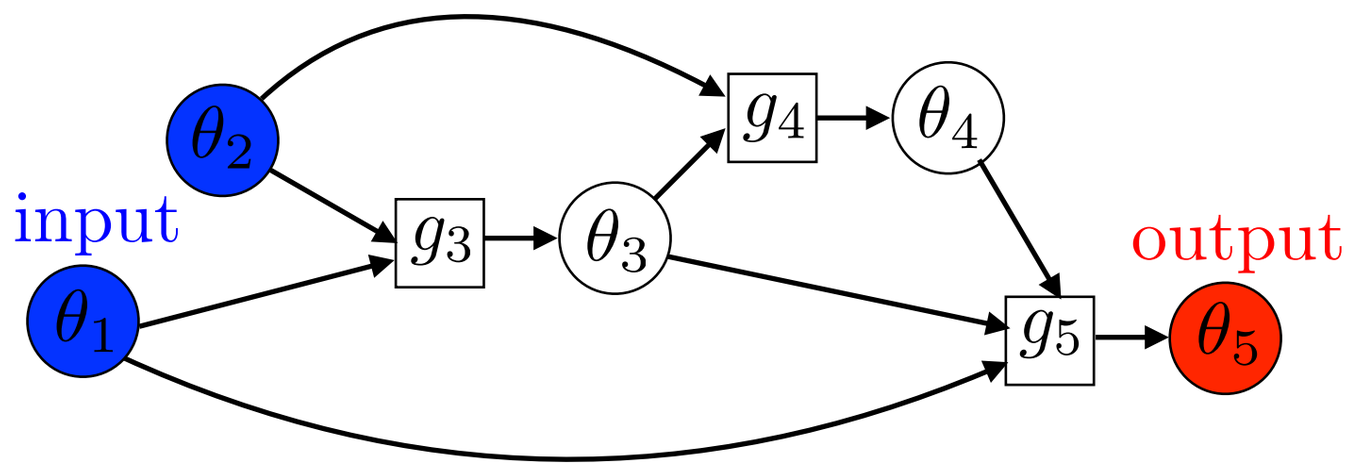

computing $\nabla \ell$

$$\text{“}\frac{\partial \ell}{\partial \theta_r} = \sum_{s \in \text{Child}(r)} \frac{\partial \ell}{\partial \theta_s}\frac{\partial \theta_s}{\partial \theta_r}\text{”}$$

$\nabla_r \ell(\theta)$ $\qquad$ $\nabla_s \ell(\theta)$ $\qquad$ $\partial_r g_s(\theta)$

What's next

Alexandre Allauzen: deep neural networks training.

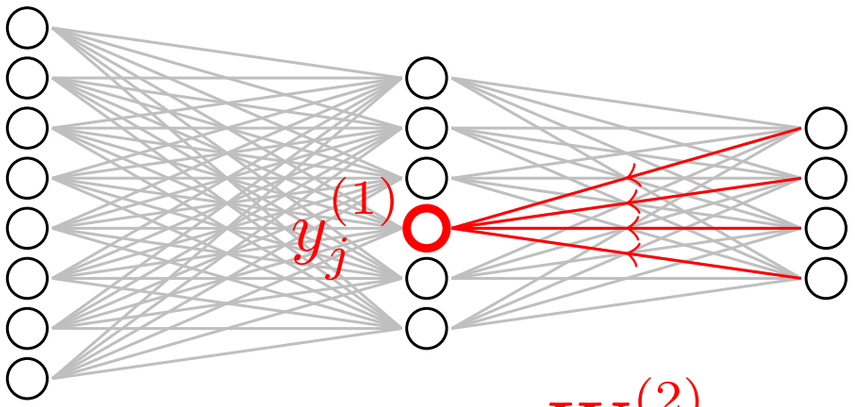

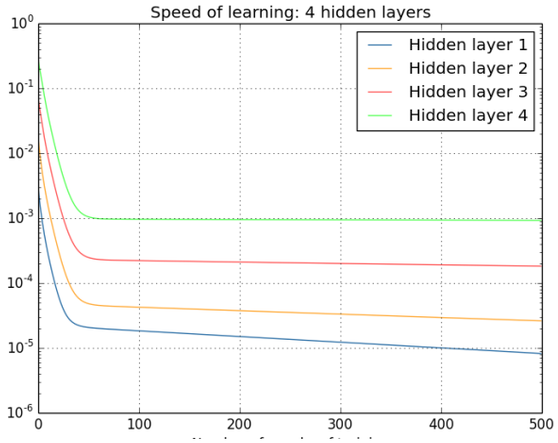

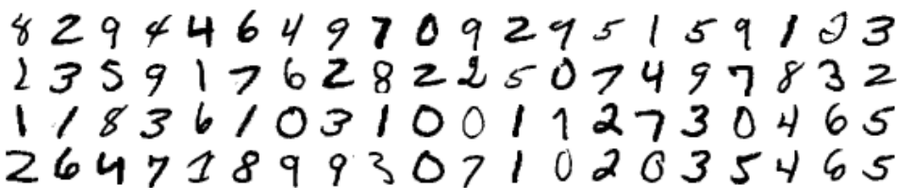

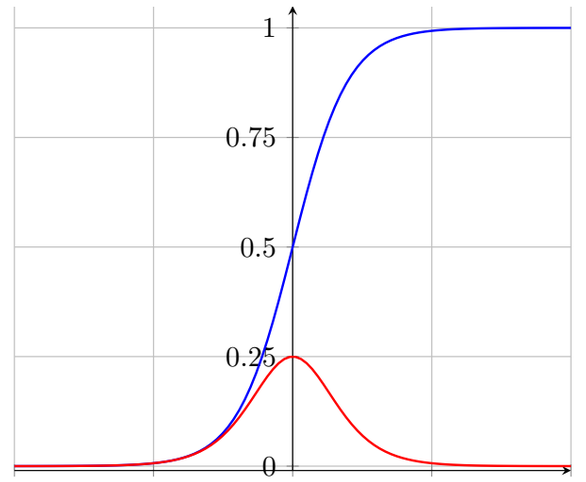

Guillaume Charpiat: architecture of deep neural networks.